

Visualizing the alignment stats of STAR using R

This report was generated with R and the help of the knitr package.

The data set used here contains mRNA from 48 replicates of two *S. cerevisiae* populations: wildtype and *snf2* knock-out mutants. All 96 samples were sequenced on one flow cell (Illumina HiSeq 2000) The SRA accession number for the entire data set (consisting of 2 x 48 samples) is ERP004763.

```
library(knitr) # for turning the original script into a nicely formatted pdf
opts_chunk$set(echo = TRUE, message = FALSE) # knitr options
opts_knit$set(self.contained = FALSE) # knitr options
library(ggplot2)
library(gridExtra) # for composite plotting of ggplots
options(stringsAsFactors = FALSE)
```

Reading in data

```
# listing files to be read in
infiles <- list.files(path="01_STAR_logs/", # folder with log files
                     pattern="Log.final.out",
                     full.names = TRUE)
```

```
# iterating over the file list generates a list of data frames
align.results <- lapply(infiles, function(x)
  read.table(x, sep="|", strip.white=TRUE, stringsAsFactor=FALSE,
            skip=3, fill = TRUE, header = FALSE))
typeof(align.results)
```

```
## [1] "list"
```

```
head(align.results[[1]])
```

```
##           V1           V2
## 1 Mapping speed, Million of reads per hour 1038.94
## 2           Number of input reads 12121025
## 3           Average input read length      51
## 4           UNIQUE READS:
## 5           Uniquely mapped reads number 10850859
## 6           Uniquely mapped reads %    89.52%
```

```
# removing "%" from some of the values to keep just the numeric parts
align.results <- lapply(align.results, function(x)
  transform(x, V2 = as.numeric(gsub("%", "", x$V2) )))

# some cosmetics of each data frame's name, specific for the sample names
# of the files used here
names(align.results) <- gsub(".*(SNF2|WT)(\\_[0-9]*).*", "\\1\\2", infiles)
```

Generating a long data frame for ggplot2

```
# concatenating all data frames of align.results
align.results.df <- as.data.frame(do.call(rbind, align.results))

# removing lines without values
align.results.df <- align.results.df[complete.cases(align.results.df),]

# adding additional columns with information about sample and replicate ID,
# using the information from the row names (which are, in turn, based on the
# names of the individual data frames that were stored in the original
# list, align.results)
align.results.df$sample <- gsub("(.*)\\"_.*", "\\1", row.names(align.results.df))
align.results.df$replicate <- as.factor(as.numeric(
  gsub(".*\\_([0-9]*)\\"_.*", "\\1", row.names(align.results.df))
))

# check the result - we should have a data frame with 4 columns
head(align.results.df)
```

```
##                               V1          V2 sample
## SNF2_1.1 Mapping speed, Million of reads per hour    1038.94   SNF2
## SNF2_1.2                Number of input reads 12121025.00   SNF2
## SNF2_1.3                Average input read length      51.00   SNF2
## SNF2_1.5                Uniquely mapped reads number 10850859.00   SNF2
## SNF2_1.6                Uniquely mapped reads %       89.52   SNF2
## SNF2_1.7                Average mapped length      50.72   SNF2
##          replicate
## SNF2_1.1          1
## SNF2_1.2          1
## SNF2_1.3          1
## SNF2_1.5          1
## SNF2_1.6          1
## SNF2_1.7          1
```

Making plots

Numbers are usually more easily digestible and comparable if they are represented visually, especially if one has many samples (although you will probably rarely encounter 48 samples per replicate as for this example data used here). However, we don't need to visualize every entry from the STAR output since some are redundant, and some even not applicable in our case:

```
unique(align.results.df$V1)
```

```
## [1] "Mapping speed, Million of reads per hour"
## [2] "Number of input reads"
## [3] "Average input read length"
## [4] "Uniquely mapped reads number"
## [5] "Uniquely mapped reads %"
## [6] "Average mapped length"
## [7] "Number of splices: Total"
```

```
## [8] "Number of splices: Annotated (sjdb)"
## [9] "Number of splices: GT/AG"
## [10] "Number of splices: GC/AG"
## [11] "Number of splices: AT/AC"
## [12] "Number of splices: Non-canonical"
## [13] "Mismatch rate per base, %"
## [14] "Deletion rate per base"
## [15] "Deletion average length"
## [16] "Insertion rate per base"
## [17] "Insertion average length"
## [18] "Number of reads mapped to multiple loci"
## [19] "% of reads mapped to multiple loci"
## [20] "Number of reads mapped to too many loci"
## [21] "% of reads mapped to too many loci"
## [22] "% of reads unmapped: too many mismatches"
## [23] "% of reads unmapped: too short"
## [24] "% of reads unmapped: other"
```

Thus, we first define those entries that we are interested in:

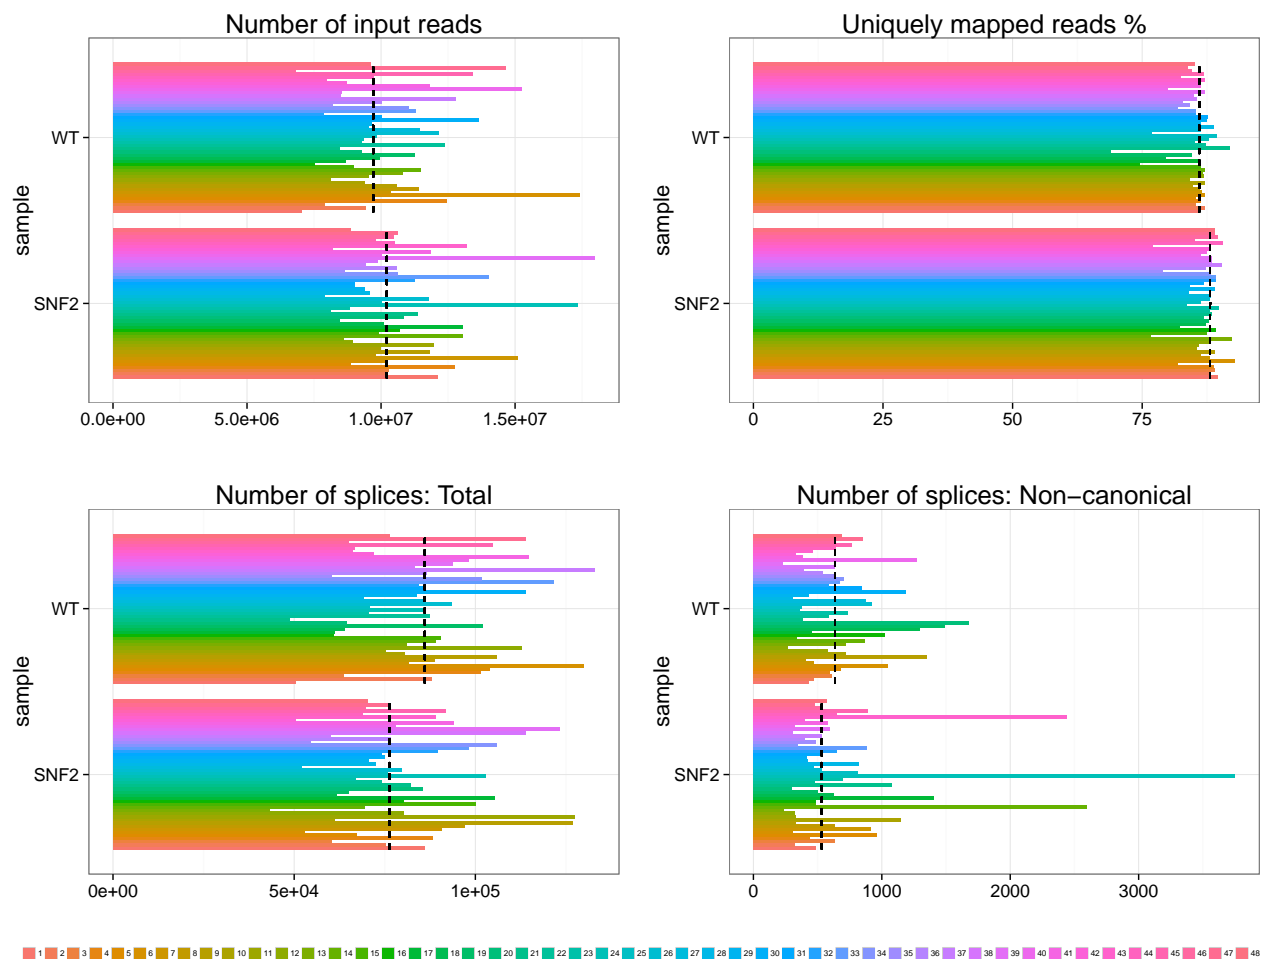
```
filters = c("Number of input reads", "Uniquely mapped reads %",
            "Number of splices: Total", "Number of splices: Non-canonical")
```

Now, let's generate some bar plots. The following chunk of code is optimized for returning a composite image with four plots and one legend. Note that the functions `PlottingAlignmentResults` and `ExtractLegend` are part of the script that we sourced in the beginning.

```
# for each entry in "filters", generate a bar chart
plots <- lapply(filters, function(x)
  PlottingAlignmentResults(x, align.results.df, Legend = FALSE))





# getting the legend is a bit tricky (kudos to Luce for figuring it out)
my.legend <- ExtractLegend(PlottingAlignmentResults(align.results.df,
  Filter = filters[1],
  Legend=TRUE))

# combining plots and legend
grid.arrange(arrangeGrob(plots[[1]], plots[[2]], plots[[3]], plots[[4]], nrow=2),
  my.legend$legend, nrow=2,
  heights= unit.c(unit(1, "npc") - my.legend$lheight, my.legend$lheight)
)
```

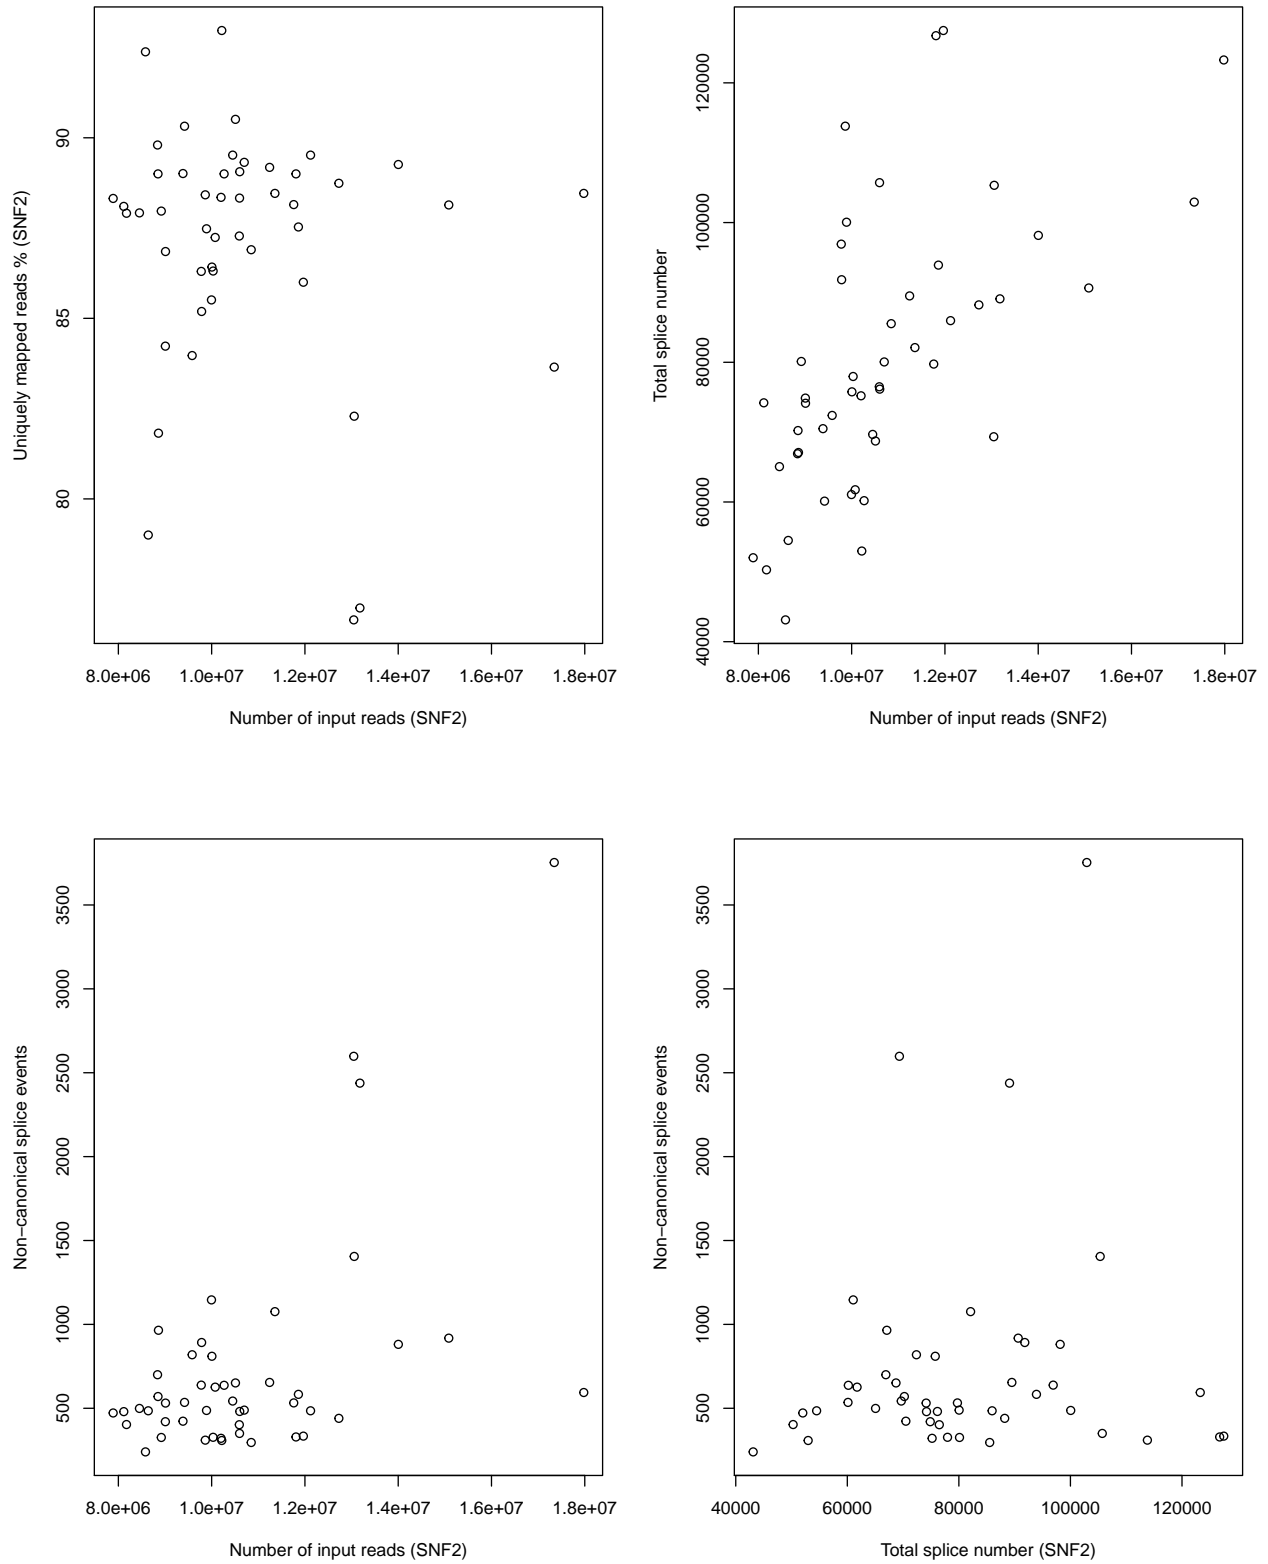


The dashed black line is the median value across all replicates of one condition. There is quite a bit of variation for all four stats between the replicates. Although the median sequencing depth of the SNF2 samples is slightly higher than for the WT samples, the median number of splice sites shows a reverse trend. This is counter-intuitive as the number of splice sites should correlate positively with the number of mapped reads. Let's check whether that relationship is simply confounded by just looking at the median values.

WT



SNF2



In both sample types, the number of detected splice sites indeed seems to correlate with the number of input reads.

Apart from that, the samples almost all aligned well (70-80% alignment rate is typically seen for RNA-seq experiments).

For future references, let's get those samples that seem to deviate the most:

- samples with low mapping rate (<75%)

```
subset(aligned.results.df, V1 == filters[2] & V2 < 75)
```

```
##                               V1    V2 sample replicate
## WT_16.6 Uniquely mapped reads % 74.64    WT           16
## WT_20.6 Uniquely mapped reads % 68.98    WT           20
```

- samples with high numbers of non-canonical splice events (>1,000)

```
subset(aligned.results.df, V1 == filters[4] & V2 > 1000)
```

```
##                               V1    V2 sample replicate
## SNF2_10.13 Number of splices: Non-canonical 1146    SNF2           10
## SNF2_14.13 Number of splices: Non-canonical 2598    SNF2           14
## SNF2_17.13 Number of splices: Non-canonical 1405    SNF2           17
## SNF2_21.13 Number of splices: Non-canonical 1076    SNF2           21
## SNF2_24.13 Number of splices: Non-canonical 3753    SNF2           24
## SNF2_43.13 Number of splices: Non-canonical 2438    SNF2           43
## WT_16.13   Number of splices: Non-canonical 1023     WT            16
## WT_18.13   Number of splices: Non-canonical 1296     WT            18
## WT_19.13   Number of splices: Non-canonical 1488     WT            19
## WT_20.13   Number of splices: Non-canonical 1679     WT            20
## WT_30.13   Number of splices: Non-canonical 1188     WT            30
## WT_40.13   Number of splices: Non-canonical 1270     WT            40
## WT_6.13    Number of splices: Non-canonical 1045     WT             6
## WT_9.13    Number of splices: Non-canonical 1351     WT             9
```